Visualizing Linear Models:
An R Bag of Tricks
Session 1: Getting Started

Michael Friendly
SCS Short Course
Oct. 22, 29, Nov. 5, 2020
https://friendly.github.io/VisMLM-course/

Today's topics

- What you need for this course
- Why plot your data?
- Data plots
- Model (effect) plots
- Diagnostic plots

What you need

- R, version >=3.6
  - Download from https://cran.r-project.org/
- RStudio IDE, highly recommended
  - https://www.rstudio.com/products/rstudio/
- R packages: see course web page
  - car
  - effects
  - heplots
  - candisc

Why plot your data?

Getting information from a table is like extracting sunlight from a cucumber. --- Farquhar & Farquhar, 1891

Information that is imperfectly acquired, is generally as imperfectly retained; and a man who has carefully investigated a printed table, finds, when done, that he has only a very faint and partial idea of what he has read; and that like a figure imprinted on sand, is soon totally erased and defaced.
--- William Playfair, *The Commercial and Political Atlas* (p. 3), 1786

Cucumbers

Table 7
Stevens et al. 2006, table 2: Determinants of authoritarian aggression

Variable	Coefficient (Standard Error)
Constant	.41 (.93)
Countries	
Argentina	1.31 (.33)**B,M
Chile	.93 (.32)**B,M
Colombia	1.46 (.32)**B,M
Mexico	.07 (.32)A,CH,CO,V
Venezuela	.96 (.37)**B,M
Threat	
Retrospective egocentric economic perceptions	.20 (.13)
Prospective egocentric economic perceptions	.22 (.12)#
Retrospective sociotropic economic perceptions	−.21 (.12)#
Prospective sociotropic economic perceptions	−.32 (.12)*
Ideological distance from president	−.27 (.07)**
Ideology	
Ideology	.23 (.07)**
Individual Differences	
Age	.00 (.01)
Female	−.03 (.21)
Education	.13 (.14)
Academic Sector	.15 (.29)
Business Sector	.31 (.25)
Government Sector	−.10 (.27)
R^2	.15
Adjusted R^2	.12
N	500

Results of a one model for authoritarian aggression

The information is overwhelmed by footnotes & significance **stars**

**p < .01, *p < .05, #p < .10 (twotailed)
A Coefficient is significantly different from Argentina's at p < .05;
B Coefficient is significantly different from Brazil's at p < .05;
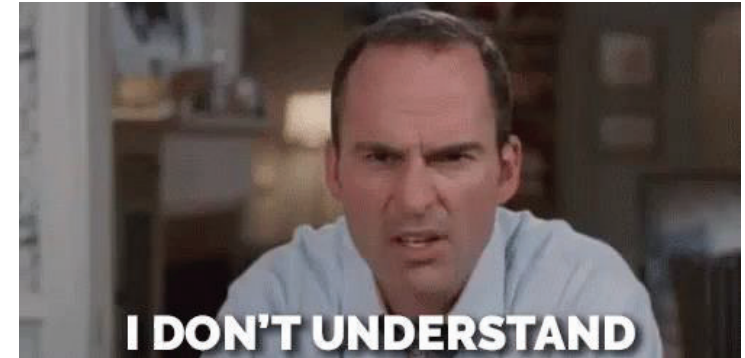CH Coefficient is significantly different from Chile's at p < .05;
CO Coefficient is significantly different from Colombia's at p < .05;
M Coefficient is significantly different from Mexico's at p < .05;
V Coefficient is significantly different from Venezuela's at p < .05.
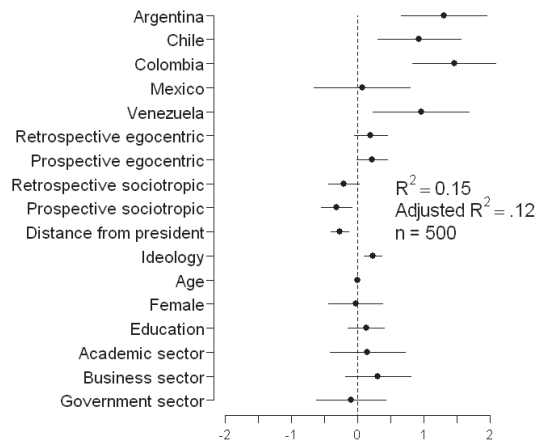
What's wrong with this picture?

I DON'T UNDERSTAND

Sunlight

coefplot(model)

$R^2 = 0.15$
Adjusted $R^2 = .12$
n = 500

Why didn't they say this in the first place?

NB: This is a presentation graph equivalent of the table

Shows coefficient with 95% CI

Run, don't walk toward the sunlight

Graphs can give enlightenment

The greatest value of a picture is when it forces us to notice what we never expected to see.
-- John W. Tukey

Effect of one rotten point on regression

Self-Reports of Height and Weight

Dangers of numbers-only output

Student: You said to run descriptives and compute the correlation. What next?

Consultant: Did you plot your data?

With exactly the same stats, the data could be *any* of these plots

See how this in done in R: https://cran.r-project.org/web/packages/datasauRus/

```
X Mean: 54.26
Y Mean: 47.83
X SD  : 16.76
Y SD  : 26.93
Corr. : -0.06
```

Sometimes, don't need numbers at all

COVID transmission risk ~ Occupancy * Ventilation * Activity * Mask? * Contact.time

A complex 5-way table, whose message is clearly shown w/o numbers

There are 1+ unusual cells here. Can you see them?

From: N.R. Jones et-al (2020). Two metres or one: what is the evidence for physical distancing in covid-19? *BMJ* 2020;370:m3223, *doi: https://doi.org/10.1136/bmj.m3223*

If you do need tables– make them pretty

Several R packages make it easier to construct informative & pretty semi-graphic tables

Presentation graph

Perhaps too cute!

Distribution of variables shown

Flipper lengths (mm) of the famous penguins of Palmer Station, Antarctica.

Species	Distribution	Female Avg.	Female Std. Dev.	Male Avg.	Male Std. Dev.
ADELIE		188	5.6	192	6.6
CHINSTRAP		192	5.8	200	6.0
GENTOO		213	3.9	222	5.7

Artwork by @allison_horst

Visual table ideas: Heatmap shading

Heatmap shading: Shade the background of each cell according to some criterion

The trends in the US and Canada are made obvious

NB: Table rows are sorted by Jan. value, lending coherence

Unemployment rate in selected countries
January-August 2020, sorted by the unemployment rate in January.

country	Jan ▲	Feb	Mar	Apr	May	Jun	Jul	Aug
Japan	2.4%	2.4%	2.5%	2.6%	2.9%	2.8%	2.9%	3.0%
Netherlands	3.0%	2.9%	2.9%	3.4%	3.6%	4.3%	4.5%	4.6%
Germany	3.4%	3.6%	3.8%	4.0%	4.2%	4.3%	4.4%	4.4%
Mexico	3.6%	3.6%	3.2%	4.8%	4.3%	5.4%	5.2%	5.0%
US	3.6%	3.5%	4.4%	14.7%	13.3%	11.1%	10.2%	8.4%
South Korea	4.0%	3.3%	3.8%	3.8%	4.5%	4.3%	4.2%	3.2%
Denmark	4.9%	4.9%	4.8%	4.9%	5.5%	6.0%	6.3%	6.1%
Belgium	5.1%	5.0%	5.0%	5.1%	5.0%	5.0%	5.0%	5.1%
Australia	5.3%	5.1%	5.2%	6.4%	7.1%	7.4%	7.5%	6.8%
Canada	5.5%	5.6%	7.8%	13.0%	13.7%	12.3%	10.9%	10.2%
Finland	6.8%	6.9%	7.0%	7.3%	7.5%	7.8%	8.0%	8.1%

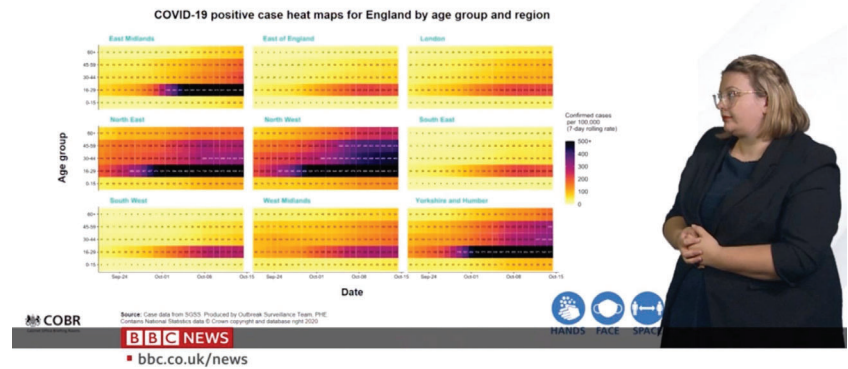Source: OECD · Get the data · Created with Datawrapper

13

Visual table ideas: Heatmap shading

As seen on TV …
Covid rate ~ Age x Date x UK region

Better: incorporate geography, not just arrange regions alphabetically

14

Visual table ideas: Sparklines

Sparklines: Mini graphics inserted into table cells or text

From: https://www.pluralsight.com/guides/tableau-playbook-sparklines

15

Linear models

- Model:

$$y_i = \beta_0 + \beta_1 X_{i1} + \beta_2 X_{i2} + \cdots + \beta_p X_{ip} + \varepsilon_i$$

- Xs: quantitative predictors, factors, interactions, …

- Assumptions:
  - **Linearity**: Predictors (possibly transformed) are linearly related to the outcome, **y**. [This just means linear in the **parameters**.]
  - **Specification**: No important predictors have been omitted; only important ones included. [This is often key & overlooked.]
  - The "holy trinity":
    - **Independence**: the errors are uncorrelated
    - **Homogeneity of variance**: Var(ε_i) = σ^2 = constant
    - **Normality**: ε_i have a normal distribution

$$\varepsilon_i \sim_{iid} \mathcal{N}(0, \sigma^2)$$

16

Plots for linear models

- Data plots:
  - plot response (y) vs. predictors, with smooth summaries
  - scatterplot matrix --- all pairs
- Model (effect) plots
  - plot predicted response ($\hat{y}$) vs. predictors, controlling for variables not shown.
- Diagnostic plots

Occupational Prestige data

- Data on prestige of 102 occupations and
  - average education (years)
  - average income ($)
  - % women
  - type (Blue Collar, Professional, White Collar)

```
> head(Prestige)
                     education income women prestige census type
gov.administrators       13.11  12351 11.16     68.8   1113 prof
general.managers         12.26  25879  4.02     69.1   1130 prof
accountants              12.77   9271 15.70     63.4   1171 prof
purchasing.officers      11.42   8865  9.11     56.8   1175 prof
chemists                 14.62   8403 11.68     73.5   2111 prof
physicists               15.64  11030  5.13     77.6   2113 prof
```

Informative scatterplots

Scatterplots are most useful when enhanced with annotations & statistical summaries

Data ellipse and regression line show the linear model, prestige ~ income

Point labels show possible outliers

Smoothed (loess) curve and CI show the trend

Boxplots show marginal distributions

Informative scatterplots

car::scatterplot() provides all these enhancements

```
scatterplot(prestige ~ income, data=Prestige,
        pch = 16,
        regLine = list(col = "red", lwd=3),
        smooth = list(smoother=loessLine,
                      lty.smooth = 1, col.smooth = "black",
                      lwd.smooth=3, col.var = "darkgreen"),
        ellipse = list(levels = 0.68),
        id = list(n=4, col="black", cex=1.2))
```

Skewed distribution of income & non-linear relation suggest need for a transformation

Arrow rule: move on the scale of powers in direction of the bulge

Try log(income)

```
scatterplot(prestige ~ income, data=Prestige,
            log = "x",
            pch = 16,
            regLine = list(col = "red", lwd=3),
            … )
```

Income now ~ symmetric

Relation closer to linear

Stratify by type?

```
scatterplot(prestige ~ income | type, data=Prestige,
      col = c("blue", "red", "darkgreen"),
      pch = 15:17,
      legend = list(coords="bottomright"),
      smooth=list(smoother=loessLine, var=FALSE, span=1, lwd=4))
```

Formula: | type → "given type"

Different slopes: interaction of income * type

Provides another explanation of the non-linear relation

This is a new finding!

# Scatterplot matrix

```
scatterplotMatrix(~ prestige + education + income + women ,
      data=Prestige,
      regLine = list(method=lm, lty=1, lwd=2, col="black"),
      smooth=list(smoother=loessLine, spread=FALSE,
            lty.smooth=1, lwd.smooth=3, col.smooth="red"),
      ellipse=list(levels=0.68, fill.alpha=0.1))
```

prestige vs. all predictors

diagonal: univariate distributions
- income: + skewed
- %women: bimodal

off-diagonal: relations among predictors
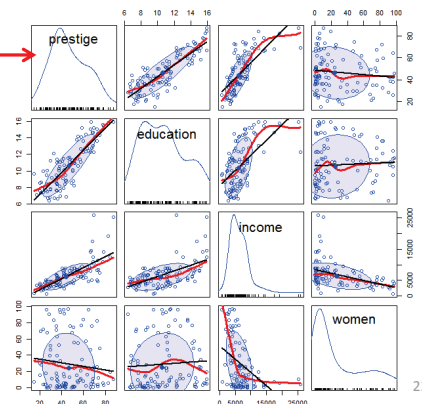
Fit a model

```
> mod1 <- lm(prestige ~ education + poly(women, 2) +
+                    log(income)*type, data=Prestige)
> summary(mod1)

Coefficients:
                    Estimate Std. Error t value Pr(>|t|)
(Intercept)         -137.500    23.522    -5.85  8.2e-08 ***
education              2.959     0.582     5.09  2.0e-06 ***
poly(women, 2)1       28.339    10.190     2.78   0.0066 **
poly(women, 2)2       12.566     7.095     1.77   0.0800 .
log(income)           17.514     2.916     6.01  4.1e-08 ***
typeprof              74.276    30.736     2.42   0.0177 *
typewc                 0.969    39.495     0.02   0.9805
log(income):typeprof  -7.698     3.451    -2.23   0.0282 *
log(income):typewc    -0.466     4.620    -0.10   0.9199
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' '
1

Multiple R-squared: 0.879,   Adjusted R-squared:  0.868
F-statistic: 81.1 on 8 and 89 DF,  p-value: <2e-16
```

- allow women^2 term
- interaction of log(income) and type

Fits very well!

Model (effect) plots

- We'd like to see plots of the predicted value ($\hat{y}$) of the response against predictors
  - But must control for other predictors not shown in a given plot
  - Variables not shown in a given plot are averaged over.
  - Slopes of lines reflect the partial coefficient in the model
  - Partial residuals can be shown also

For details, see `vignette("predictor-effects-gallery", package="effects)`

25

Model (effect) plots: education

```
library("effects")
mod1.e1 <- predictorEffect("education", mod1)
plot(mod1.e1)
```

This graph shows the partial slope for education.

For each ↑ year in education, fitted prestige ↑2.96 points, (other predictors held fixed)

education predictor effect plot

---

Model (effect) plots

```
mod1.e1a <- predictorEffect("education", mod1, residuals=TRUE)
plot(mod1.e1a,
     residuals.pch=16, id=list(n=4, col="black"))
```

Partial residuals show the residual of prestige controlling for other predictors

Unusual points here would signal undue influence

education predictor effect plot

27

Model (effect) plots: women

```
mod1.e2 <- predictorEffect("women", mod1, residuals=TRUE)
plot(mod1.e2, ylim=c(40, 65), lwd=4,
     residuals.pch=16)
```

Surprise!
Prestige of occupations ↑ with % women (controlling for other variables


women predictor effect plot

Model (effect) plots: income

```
plot(predictorEffect("income", mod1),
     lines=list(multiline=TRUE, lwd=3),
     key.args = list(x=.7, y=.35))
```

Income interacts with type in the model

The plot is curved because log(income) is in the model

income predictor effect plot

Diagnostic plots

- The linear model, $y = X\beta + \epsilon$ assumes:
  - Residuals, ε_i are normally distributed, $\varepsilon_i \sim N(0,\sigma^2)$
  - (Normality not required for **X**s)
  - Constant variance, $Var(\varepsilon_i) = \sigma^2$
  - Observations y_i are statistically independent
- Violations $\rightarrow$ inferences may not be valid
- A variety of plots can diagnose all these problems
- Other methods (boxCox, boxTidwell) diagnose the need for transformations of **y** or **X**s.

The "regression quartet"

In R, plotting a `lm` model object $\rightarrow$ the "regression quartet" of plots

```
plot(mod1, lwd=2, cex.lab=1.4)
```

❶ Residuals: should be flat vs. fitted values

❷ Q-Q plot: should follow the 45° line

❸ Scale-location: should be flat if constant variance

❹ Resids vs. leverage: can show influential observations

Unusual data: Leverage & Influence

- "Unusual" observations can have dramatic effects on least-squares estimates in linear models
- Three archetypal cases:
  - Typical X (low leverage), bad fit   -- Not much harm
  - Unusual X (high leverage), good fit   -- Not much harm
  - Unusual X (high leverage), bad fit   -- BAD, BAD, BAD
- Influential observations: unusual in *both* X & Y
- Heuristic formula:

  Influence = X leverage x Y residual

## Influence plots

Influence (Cook's D) measures impact of individual obs. on coefficients, fitted values

Influence ~ Residual (y -ŷ) × Hat-value $(X - \bar{X})^2$

Bubble size ~ influence

`influencePlot(mod1)`

Spread-level plots

- To diagnose non-constant variance, plot:
  - log |Std. residual| vs. log (x)
  - log (IQR) vs log (median)   [for grouped data]
- If ≈ linear w/ slope b, transform $y \rightarrow y^{(1-b)}$

Artificial data, generated so σ ~ x
- b ≈ 1 → power = 0
- → analyze log(y)

Spread-level plot: baseball data

Data on salary and batter performance from 1987 season

```
data("Baseball", package="vcd")
bb.mod <- lm(sal87 ~ years + hits + runs + homeruns, data=Baseball)
spreadLevelPlot(bb.mod, pch=16, lwd=3,
                id=list(n=2))

## Suggested power transformation: 0.2609
```

slope = .74 → p = .26

i.e., $y \rightarrow \log(y)$ or $y^{1/4}$

NB: both axes plotted on log scale

Summary

- Tables are for look-up; graphs can give insight
- Data plots are more effective when enhanced
  - data ellipses → strength & precision of correlation
  - regression lines and smoothed curves
  - point identification → noteworthy observations
- Effect plots show informative views of models
- Diagnostic plots can reveal influential observations and need for transformations.